

AAMAS System Architecture & Design

Autonomous AI Multi-Agent Software Development Platform

SYSTEM OVERVIEW

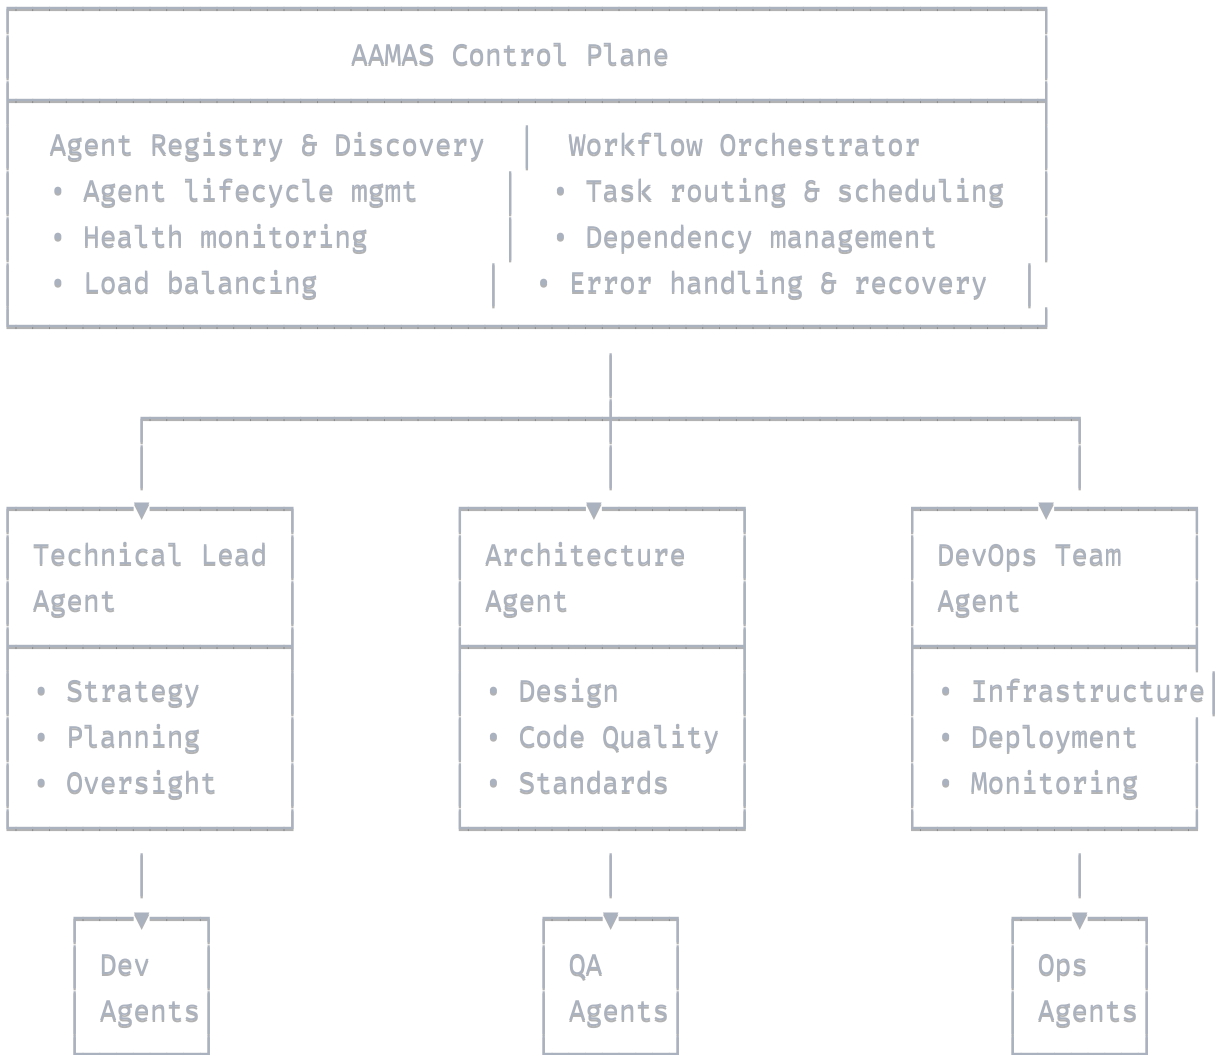
AAMAS is designed as a cloud-native, microservices-based platform that orchestrates specialized AI agents to automate software development workflows. The architecture emphasizes scalability, reliability, and intelligent agent coordination.

High-Level Architecture Principles

- **Agent-Centric Design:** Every major function is handled by specialized AI agents
 - **Event-Driven Architecture:** Asynchronous communication between agents and services
 - **Microservices Pattern:** Loosely coupled, independently deployable services
 - **Multi-Tenant SaaS:** Secure isolation between customer environments
 - **Cloud-Native:** Kubernetes-based deployment with auto-scaling capabilities
-

CORE SYSTEM ARCHITECTURE

1. Agent Orchestration Layer



2. Multi-Agent Communication System

python

Agent Communication Protocol

```
class AgentCommunicationBus:
    def __init__(self):
        self.message_broker = KafkaProducer()
        self.agent_registry = AgentRegistry()

    async def send_task(self, from_agent: str, to_agent: str, task: Task):
        """Route tasks between agents with priority and routing logic"""
        message = TaskMessage(
            sender=from_agent,
            recipient=to_agent,
            task=task,
            timestamp=datetime.utcnow(),
            correlation_id=uuid4()
        )

        # Route through appropriate channel based on task type
        topic = self.get_routing_topic(task.type)
        await self.message_broker.send(topic, message)

    async def broadcast_event(self, event: SystemEvent):
        """Broadcast system-wide events to all relevant agents"""
        interested_agents = self.agent_registry.get_subscribers(event.type)
        for agent_id in interested_agents:
            await self.send_event(agent_id, event)
```

3. Agent State Management

python

```
class AgentStateManager:
    """Manages persistent state for all agents across the system"""

    def __init__(self):
        self.state_store = RedisCluster() # For fast access
        self.persistent_store = MongoDB() # For long-term storage

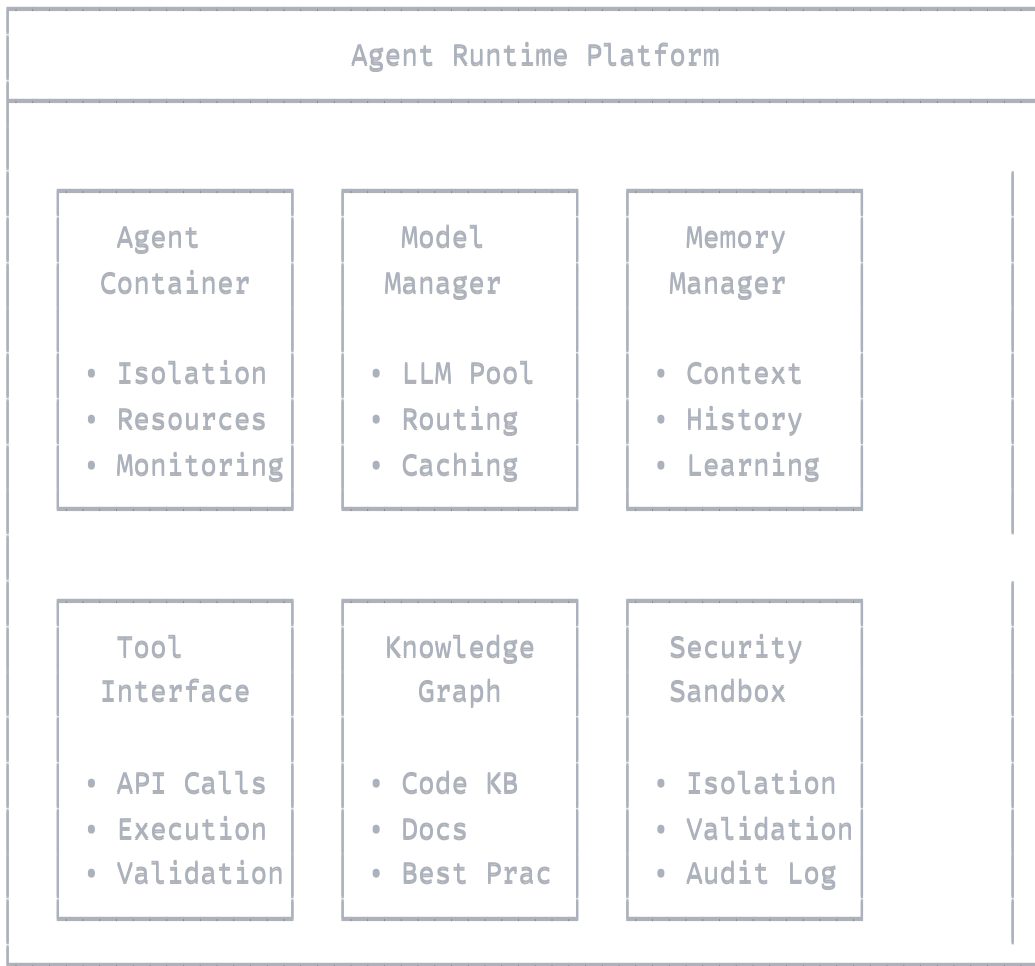
    async def save_agent_state(self, agent_id: str, state: dict):
        """Save agent state with versioning and history"""
        state_record = AgentState(
            agent_id=agent_id,
            state=state,
            version=await self.get_next_version(agent_id),
            timestamp=datetime.utcnow()
        )

        # Fast cache for active state
        await self.state_store.set(f"agent:{agent_id}:current", state_record)

        # Persistent storage for history and recovery
        await self.persistent_store.insert_one(state_record.dict())
```

DETAILED COMPONENT ARCHITECTURE

1. AI Agent Runtime Environment



2. Specialized Agent Implementations

Technical Lead Agent

python

```
class TechnicalLeadAgent(BaseAgent):
    """Strategic oversight and technical decision-making"""

    def __init__(self):
        super().__init__(
            role="technical_lead",
            capabilities=[
                "strategic_planning",
                "risk_assessment",
                "resource_allocation",
                "quality_oversight"
            ]
        )
        self.decision_engine = TechnicalDecisionEngine()
        self.risk_analyzer = RiskAnalyzer()

    async def evaluate_project_health(self, project_id: str):
        """Comprehensive project health assessment"""
        metrics = await self.get_project_metrics(project_id)
        risks = await self.risk_analyzer.analyze(project_id)
        team_performance = await self.assess_team_velocity(project_id)

        health_score = self.calculate_health_score(
            metrics, risks, team_performance
        )

        recommendations = await self.generate_recommendations(
            health_score, risks
        )

        return ProjectHealthReport(
            score=health_score,
            risks=risks,
            recommendations=recommendations,
            next_review=datetime.utcnow() + timedelta(days=7)
        )
```

Architecture Agent

python

```

class ArchitectureAgent(BaseAgent):
    """System design and code quality optimization"""

    def __init__(self):
        super().__init__(
            role="architecture",
            capabilities=[
                "system_design",
                "code_review",
                "pattern_recognition",
                "refactoring_suggestions"
            ]
        )
        self.design_patterns = DesignPatternLibrary()
        self.quality_analyzer = CodeQualityAnalyzer()

    async def review_architecture_change(self, change_request: ArchitectureChange):
        """Analyze proposed architecture changes for impact and quality"""

        # Analyze current system state
        current_architecture = await self.analyze_current_system(
            change_request.project_id
        )

        # Simulate proposed changes
        impact_analysis = await self.simulate_change_impact(
            current_architecture, change_request
        )

        # Check against best practices
        compliance_check = await self.check_design_principles(
            change_request.proposed_design
        )

        # Generate recommendations
        recommendations = await self.generate_architecture_recommendations(
            impact_analysis, compliance_check
        )

        return ArchitectureReview(
            impact_score=impact_analysis.risk_level,
            compliance_score=compliance_check.score,
            recommendations=recommendations,

```



```
        approval_status=self.determine_approval_status(impact_analysis)
    )
```

3. Development Team Agents

python

```

class DevelopmentTeamAgent(BaseAgent):
    """Feature implementation and code generation"""

    def __init__(self, specialization: str):
        super().__init__(
            role=f"developer_{specialization}",
            capabilities=[
                "code_generation",
                "testing",
                "debugging",
                "documentation"
            ]
        )
        self.code_generator = CodeGenerator(specialization)
        self.test_generator = TestGenerator()

    async def implement_feature(self, feature_spec: FeatureSpecification):
        """Full feature implementation with tests and documentation"""

        # Generate implementation plan
        implementation_plan = await self.plan_implementation(feature_spec)

        # Generate code
        code_artifacts = await self.code_generator.generate_code(
            feature_spec, implementation_plan
        )

        # Generate tests
        test_suite = await self.test_generator.generate_tests(
            code_artifacts, feature_spec
        )

        # Generate documentation
        documentation = await self.generate_documentation(
            code_artifacts, feature_spec
        )

        # Self-review
        review_result = await self.self_review_code(code_artifacts)

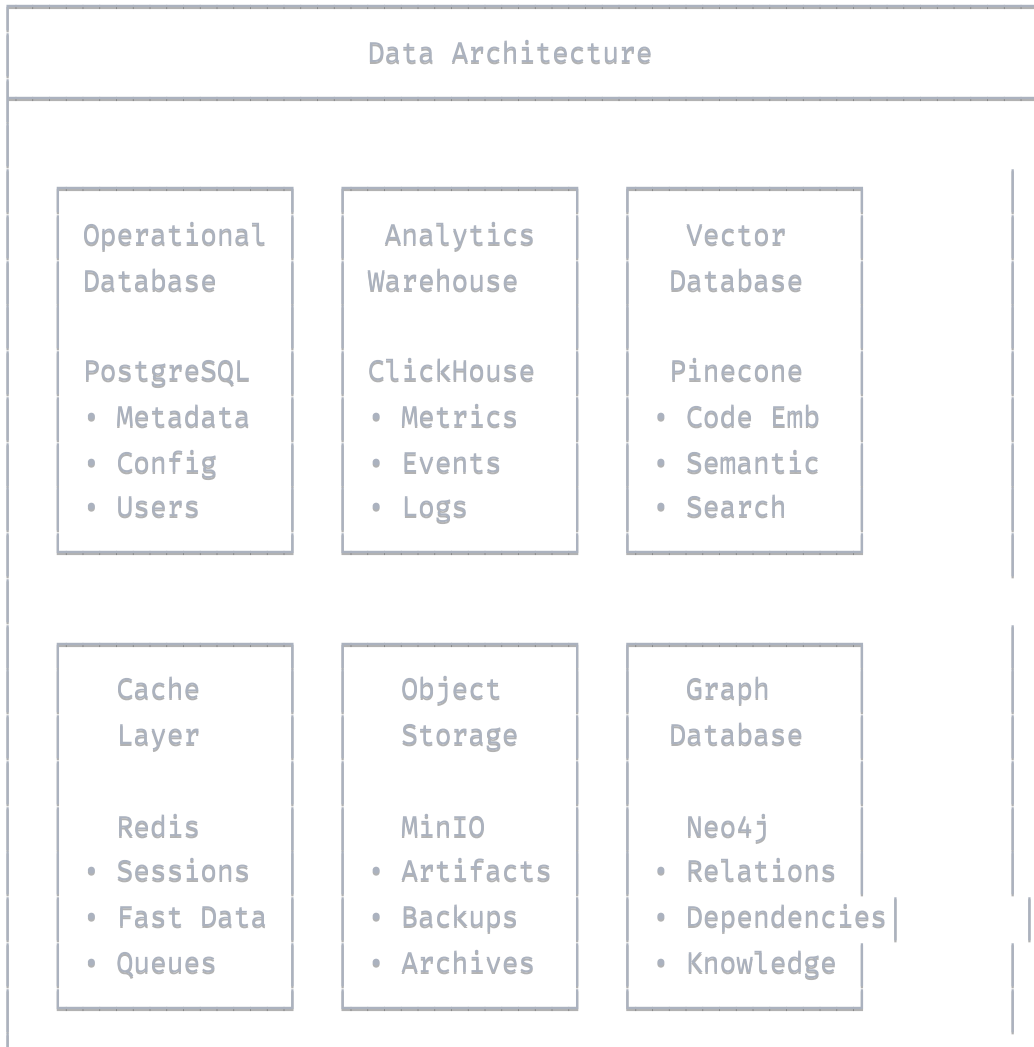
        return FeatureImplementation(
            code=code_artifacts,
            tests=test_suite,

```

```
documentation=documentation,  
quality_metrics=review_result.metrics,  
estimated_effort=implementation_plan.effort_estimate  
)
```

DATA ARCHITECTURE

1. Multi-Modal Data Storage Strategy



2. Data Models

python

```
# Core Data Models
```

```
from dataclasses import dataclass
from typing import List, Dict, Optional
from enum import Enum
```

```
@dataclass
```

```
class Project:
    id: str
    name: str
    organization_id: str
    repository_url: str
    tech_stack: List[str]
    agents_assigned: List[str]
    configuration: Dict
    created_at: datetime
    updated_at: datetime
```

```
@dataclass
```

```
class Agent:
    id: str
    type: AgentType
    specialization: str
    model_version: str
    capabilities: List[str]
    current_projects: List[str]
    performance_metrics: Dict
    status: AgentStatus
    last_active: datetime
```

```
@dataclass
```

```
class Task:
    id: str
    type: TaskType
    priority: Priority
    assigned_agent: str
    project_id: str
    description: str
    requirements: Dict
    dependencies: List[str]
    status: TaskStatus
    created_at: datetime
    due_date: Optional[datetime]
```

```
@dataclass
class CodeArtifact:
    id: str
    project_id: str
    file_path: str
    content: str
    language: str
    generated_by: str
    quality_score: float
    test_coverage: float
    documentation_score: float
    created_at: datetime
```

INTEGRATION ARCHITECTURE

1. External Tool Integration Layer

python


```

class IntegrationManager:
    """Manages all external tool integrations"""

    def __init__(self):
        self.integrations = {
            'git': GitIntegration(),
            'ci_cd': CICDIntegration(),
            'monitoring': MonitoringIntegration(),
            'communication': CommunicationIntegration()
        }

    async def register_webhook(self, tool: str, project_id: str, events: List[str]):
        """Register webhooks for external tool events"""
        integration = self.integrations[tool]
        webhook_url = f"{self.base_url}/webhooks/{tool}/{project_id}"

        await integration.register_webhook(webhook_url, events)

    async def handle_webhook(self, tool: str, project_id: str, payload: dict):
        """Process incoming webhook events"""
        event = self.parse_webhook_event(tool, payload)

        # Route to appropriate agents
        interested_agents = await self.get_interested_agents(event)

        for agent_id in interested_agents:
            await self.notify_agent(agent_id, event)

class GitIntegration:
    """Git repository integration"""

    async def setup_project_integration(self, project: Project):
        """Set up git hooks and monitoring for a project"""

        # Set up branch protection rules
        await self.setup_branch_protection(project.repository_url)

        # Configure automated PR creation
        await self.setup_pr_automation(project)

        # Set up commit analysis
        await self.setup_commit_hooks(project)

```

```

async def create_pull_request(self, project_id: str, changes: CodeChanges):
    """Create PR with AI-generated description and reviews"""

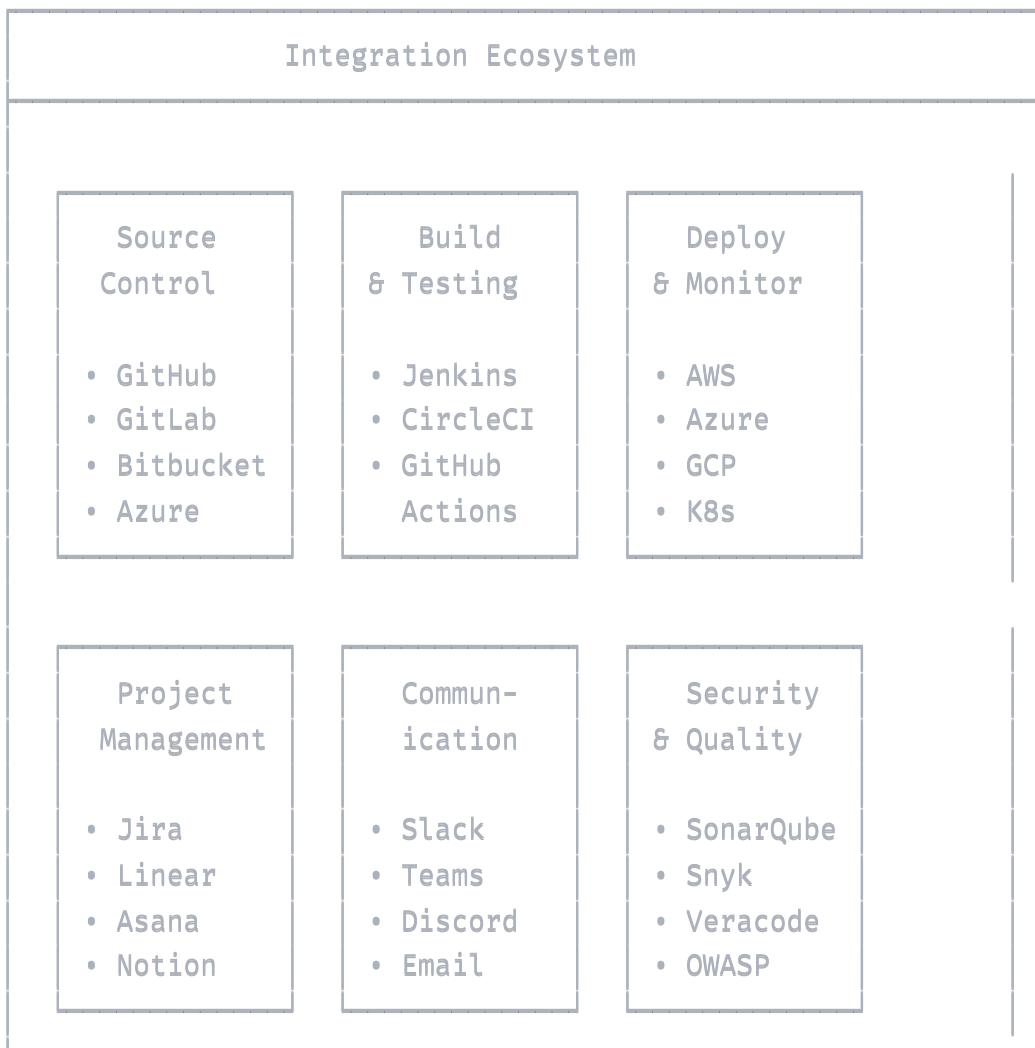
    pr_description = await self.generate_pr_description(changes)
    reviewers = await self.suggest_reviewers(changes)

    pr = await self.git_client.create_pr(
        title=changes.title,
        description=pr_description,
        reviewers=reviewers,
        labels=changes.labels
    )

    return pr

```

2. Development Tool Ecosystem



SECURITY ARCHITECTURE

1. Multi-Layered Security Model

python

```

class SecurityManager:
    """Comprehensive security management for the platform"""

    def __init__(self):
        self.auth_manager = AuthenticationManager()
        self.encryption_manager = EncryptionManager()
        self.audit_logger = AuditLogger()
        self.access_controller = AccessController()

    async def secure_agent_communication(self, message: AgentMessage):
        """Secure all inter-agent communications"""

        # Encrypt message content
        encrypted_content = await self.encryption_manager.encrypt(
            message.content, message.recipient
        )

        # Add digital signature
        signature = await self.create_message_signature(
            message.sender, encrypted_content
        )

        # Log communication for audit
        await self.audit_logger.log_agent_communication(
            sender=message.sender,
            recipient=message.recipient,
            message_type=message.type,
            timestamp=datetime.utcnow()
        )

        return SecureMessage(
            content=encrypted_content,
            signature=signature,
            metadata=message.metadata
        )

class CodeSecurityAnalyzer:
    """Analyze code for security vulnerabilities"""

    async def analyze_code_security(self, code: str, language: str):
        """Comprehensive security analysis of generated code"""

        findings = []

```

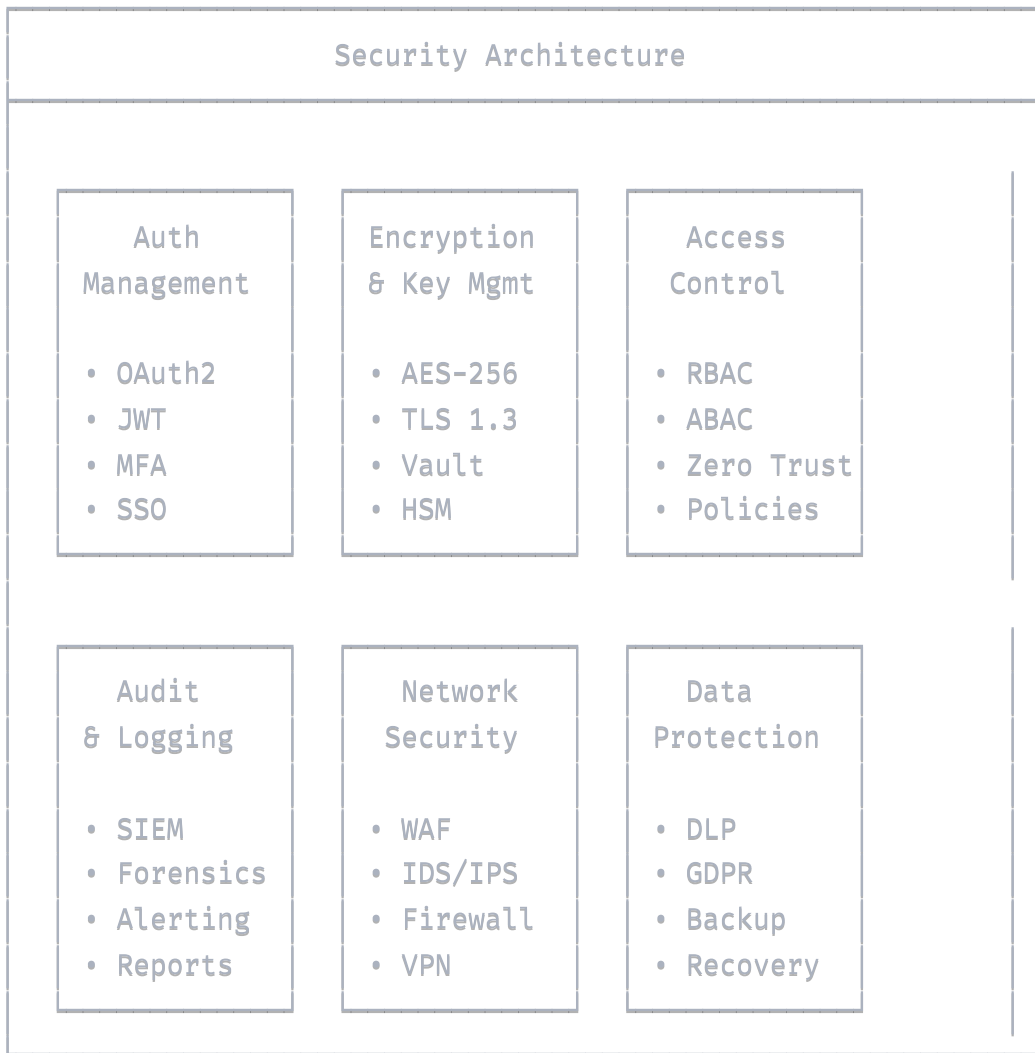
```
# Static security analysis
static_results = await self.static_security_scan(code, language)
findings.extend(static_results)

# Pattern-based vulnerability detection
pattern_results = await self.pattern_vulnerability_scan(code)
findings.extend(pattern_results)

# Dependency security check
dependency_results = await self.dependency_security_scan(code)
findings.extend(dependency_results)

return SecurityAnalysisResult(
    findings=findings,
    risk_level=self.calculate_risk_level(findings),
    recommendations=self.generate_security_recommendations(findings)
)
```

2. Data Protection & Privacy



SCALABILITY & PERFORMANCE

1. Auto-Scaling Architecture

python


```

class ScalingManager:
    """Manages dynamic scaling of agents and infrastructure"""

    def __init__(self):
        self.metrics_collector = MetricsCollector()
        self.resource_manager = ResourceManager()
        self.agent_pool = AgentPool()

    async def monitor_and_scale(self):
        """Continuous monitoring and scaling decisions"""

        while True:
            # Collect system metrics
            metrics = await self.metrics_collector.get_current_metrics()

            # Analyze scaling needs
            scaling_decision = await self.analyze_scaling_needs(metrics)

            if scaling_decision.should_scale_up:
                await self.scale_up_agents(scaling_decision.agent_types)
            elif scaling_decision.should_scale_down:
                await self.scale_down_agents(scaling_decision.agent_types)

            # Infrastructure scaling
            if scaling_decision.needs_infrastructure_scaling:
                await self.scale_infrastructure(scaling_decision.infrastructure_need)

            await asyncio.sleep(30) # Check every 30 seconds

    async def scale_up_agents(self, agent_types: List[str]):
        """Dynamically add agent instances"""

        for agent_type in agent_types:
            # Determine optimal scaling strategy
            scaling_config = await self.get_scaling_config(agent_type)

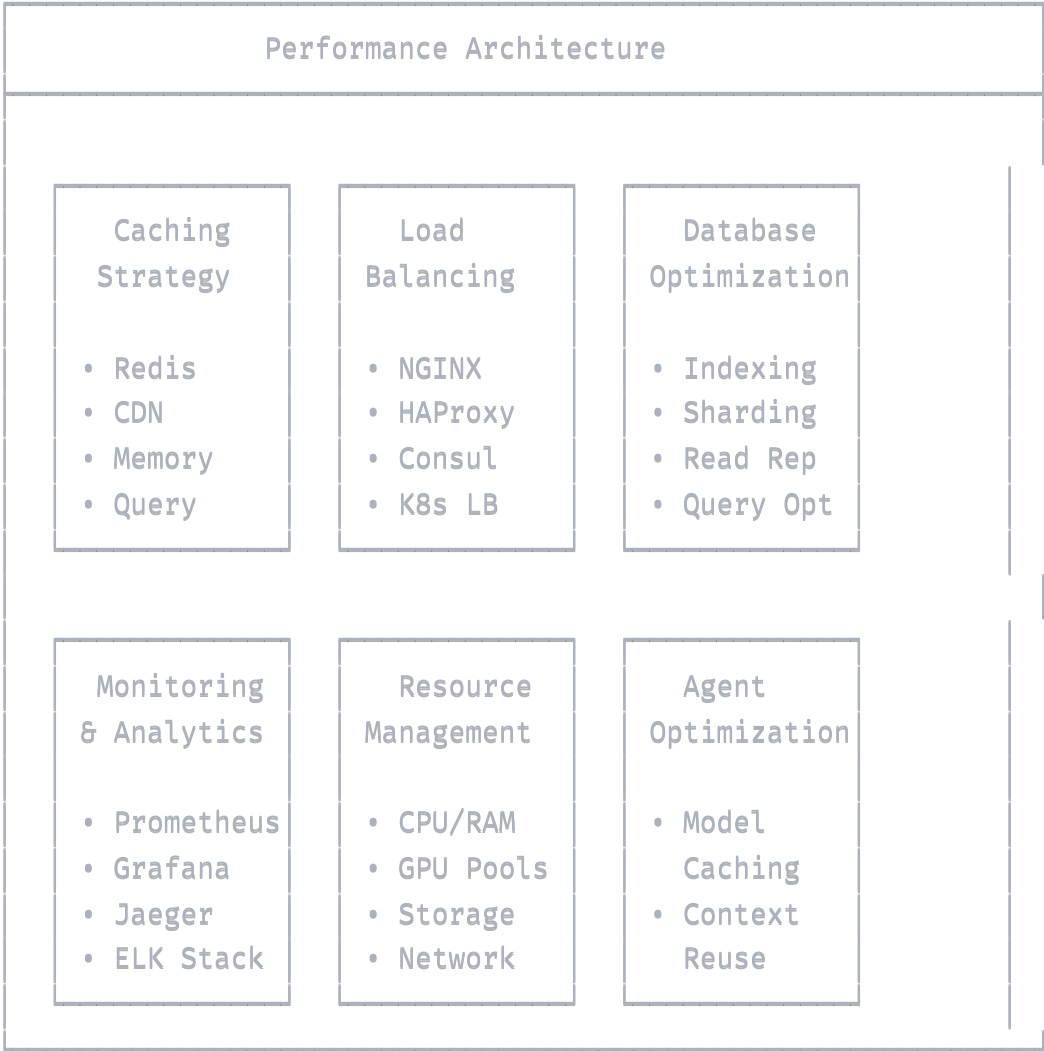
            # Create new agent instances
            new_agents = await self.agent_pool.create_agents(
                agent_type, scaling_config.instance_count
            )

            # Register and activate
            for agent in new_agents:

```

```
await self.register_agent(agent)
await self.activate_agent(agent)
```

2. Performance Optimization



DEPLOYMENT ARCHITECTURE

1. Container Orchestration

yaml

Kubernetes Deployment Configuration

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: aamas-technical-lead-agent
spec:
  replicas: 3
  selector:
    matchLabels:
      app: technical-lead-agent
  template:
    metadata:
      labels:
        app: technical-lead-agent
    spec:
      containers:
        - name: technical-lead-agent
          image: aamas/technical-lead-agent:v1.0.0
          resources:
            requests:
              memory: "2Gi"
              cpu: "1000m"
              nvidia.com/gpu: 1
            limits:
              memory: "4Gi"
              cpu: "2000m"
              nvidia.com/gpu: 1
          env:
            - name: AGENT_TYPE
              value: "technical_lead"
            - name: MODEL_ENDPOINT
              valueFrom:
                secretKeyRef:
                  name: model-config
                  key: endpoint
          ports:
            - containerPort: 8080
          livenessProbe:
            httpGet:
              path: /health
              port: 8080
            initialDelaySeconds: 30
            periodSeconds: 10
```

```
readinessProbe:  
  httpGet:  
    path: /ready  
    port: 8080  
  initialDelaySeconds: 5  
  periodSeconds: 5
```

2. CI/CD Pipeline

python

```
# CI/CD Pipeline Configuration
```

```
class AAMASPipeline:
    """Automated build, test, and deployment pipeline"""

    stages = [
        "code_quality_check",
        "security_scan",
        "unit_tests",
        "integration_tests",
        "agent_performance_tests",
        "staging_deployment",
        "end_to_end_tests",
        "production_deployment",
        "monitoring_setup"
    ]

    async def execute_pipeline(self, commit: Commit):
        """Execute full CI/CD pipeline"""

        pipeline_run = PipelineRun(
            commit=commit,
            stages=self.stages,
            started_at=datetime.utcnow()
        )

        for stage in self.stages:
            stage_result = await self.execute_stage(stage, commit)
            pipeline_run.add_stage_result(stage_result)

            if not stage_result.success:
                await self.handle_pipeline_failure(pipeline_run, stage)
                break

        return pipeline_run
```

MONITORING & OBSERVABILITY

1. Comprehensive Monitoring Stack

python

```

class MonitoringSystem:
    """Complete observability for the AAMAS platform"""

    def __init__(self):
        self.metrics_collector = PrometheusCollector()
        self.log_aggregator = ELKStack()
        self.tracing_system = Jaeger()
        self.alerting_system = AlertManager()

    async def setup_agent_monitoring(self, agent: Agent):
        """Set up comprehensive monitoring for an agent"""

        # Performance metrics
        await self.metrics_collector.register_agent_metrics(
            agent_id=agent.id,
            metrics=[
                "response_time",
                "task_completion_rate",
                "error_rate",
                "resource_utilization",
                "model_accuracy"
            ]
        )

        # Distributed tracing
        await self.tracing_system.setup_agent_tracing(agent.id)

        # Log aggregation
        await self.log_aggregator.configure_agent_logging(
            agent_id=agent.id,
            log_level="INFO"
        )

        # Alerts
        await self.setup_agent_alerts(agent)

    async def setup_agent_alerts(self, agent: Agent):
        """Configure alerting for agent health and performance"""

        alerts = [
            Alert(
                name=f"{agent.id}_high_response_time",
                condition="avg_response_time > 5s",

```



```

        severity="warning",
        actions=["scale_up", "notify_ops"]
    ),
    Alert(
        name=f"{agent.id}_error_rate",
        condition="error_rate > 5%",
        severity="critical",
        actions=["restart_agent", "notify_ops", "fallback_mode"]
    )
]

```

```

for alert in alerts:
    await self.alerting_system.register_alert(alert)

```

2. Real-Time Dashboard & Analytics

```

```python
class AAMASDashboard:
 """Real-time monitoring dashboard for AAMAS platform"""

 def __init__(self):
 self.dashboard_server = GrafanaDashboard()
 self.real_time_updater = WebSocketManager()
 self.analytics_engine = AnalyticsEngine()

 async def create_customer_dashboard(self, customer_id: str):
 """Create customized dashboard for each customer"""

 dashboard_config = {
 "customer_id": customer_id,
 "panels": [
 {
 "title": "Agent Performance Overview",
 "type": "stat_panel",
 "metrics": [
 "active_agents",
 "tasks_completed_today",
 "average_response_time",
 "success_rate"
]
 },
 {
 "title": "Development Velocity",
 "type": "time_series",

```

```

 "metrics": [
 "commits_per_day",
 "pull_requests_created",
 "bugs_fixed",
 "features_delivered"
]
 },
 {
 "title": "Code Quality Trends",
 "type": "gauge",
 "metrics": [
 "code_quality_score",
 "test_coverage",
 "technical_debt_ratio",
 "security_score"
]
 },
 {
 "title": "Resource Utilization",
 "type": "heatmap",
 "metrics": [
 "cpu_usage",
 "memory_usage",
 "gpu_utilization",
 "storage_usage"
]
 }
]

}

dashboard = await self.dashboard_server.create_dashboard(dashboard_config)

Set up real-time updates
await self.real_time_updater.setup_dashboard_updates(
 customer_id, dashboard.id
)

return dashboard

class AnalyticsEngine:
 """Advanced analytics for performance optimization"""

 async def generate_insights(self, customer_id: str, timeframe: str):
 """Generate actionable insights from platform usage"""

```

```
Collect data
usage_data = await self.collect_usage_data(customer_id, timeframe)
performance_data = await self.collect_performance_data(customer_id, timeframe)
quality_data = await self.collect_quality_data(customer_id, timeframe)

AI-powered analysis
insights = await self.ai_analyzer.analyze_patterns(
 usage_data, performance_data, quality_data
)

Generate recommendations
recommendations = await self.generate_recommendations(insights)

return AnalyticsReport(
 timeframe=timeframe,
 insights=insights,
 recommendations=recommendations,
 roi_metrics=self.calculate_roi_metrics(usage_data, performance_data)
)
```

---

## API ARCHITECTURE

### 1. RESTful API Design

python

```

from fastapi import FastAPI, HTTPSRedirect, Depends
from fastapi.security import HTTPBearer
from fastapi.middleware.cors import CORSMiddleware

app = FastAPI(
 title="AAMAS Platform API",
 description="Autonomous AI Multi-Agent Software Development Platform",
 version="1.0.0"
)

Security middleware
security = HTTPBearer()
app.add_middleware(HTTPSRedirect)
app.add_middleware(
 CORSMiddleware,
 allow_origins=["https://*.aamas.dev"],
 allow_credentials=True,
 allow_methods=["*"],
 allow_headers=["*"],
)

@app.post("/api/v1/projects", response_model=Project)
async def create_project(
 project_data: CreateProjectRequest,
 current_user: User = Depends(get_current_user)
):
 """Create a new project with AI agent assignment"""

 # Validate project configuration
 validation_result = await validate_project_config(project_data)
 if not validation_result.is_valid:
 raise HTTPException(400, validation_result.errors)

 # Create project
 project = await project_service.create_project(
 project_data, current_user.organization_id
)

 # Auto-assign appropriate agents
 agent_assignments = await agent_manager.assign_agents_to_project(
 project.id, project_data.requirements
)

```

```

 # Initialize project environment
 await project_initializer.setup_project_environment(
 project, agent_assignments
)

 return project

@app.get("/api/v1/projects/{project_id}/agents", response_model=List[AgentStatus])
async def get_project_agents(
 project_id: str,
 current_user: User = Depends(get_current_user)
):
 """Get all agents assigned to a project with their current status"""

 # Check permissions
 await check_project_access(project_id, current_user)

 # Get agent assignments
 agents = await agent_manager.get_project_agents(project_id)

 # Get current status for each agent
 agent_statuses = []
 for agent in agents:
 status = await agent_monitor.get_agent_status(agent.id)
 agent_statuses.append(status)

 return agent_statuses

@app.post("/api/v1/projects/{project_id}/tasks", response_model=Task)
async def create_task(
 project_id: str,
 task_data: CreateTaskRequest,
 current_user: User = Depends(get_current_user)
):
 """Create a new task and assign to appropriate agent"""

 await check_project_access(project_id, current_user)

 # Analyze task requirements
 task_analysis = await task_analyzer.analyze_task(task_data)

 # Find best agent for the task
 optimal_agent = await agent_matcher.find_optimal_agent(
 project_id, task_analysis
)

```

```

)

 # Create and assign task
 task = await task_service.create_task(
 project_id=project_id,
 task_data=task_data,
 assigned_agent=optimal_agent.id,
 estimated_effort=task_analysis.estimate_effort
)

 # Notify agent
 await agent_communicator.assign_task(optimal_agent.id, task)

 return task

@app.get("/api/v1/analytics/dashboard/{customer_id}")
async def get_dashboard_data(
 customer_id: str,
 timeframe: str = "7d",
 current_user: User = Depends(get_current_user)
):
 """Get dashboard analytics data for a customer"""

 await check_customer_access(customer_id, current_user)

 # Generate real-time analytics
 dashboard_data = await analytics_engine.generate_dashboard_data(
 customer_id, timeframe
)

 return dashboard_data

WebSocket endpoint for real-time updates
@app.websocket("/ws/project/{project_id}/updates")
async def project_updates_websocket(
 websocket: WebSocket,
 project_id: str,
 token: str = Query(...)
):
 """Real-time project updates via WebSocket"""

 # Authenticate WebSocket connection
 user = await authenticate_websocket_token(token)
 await check_project_access(project_id, user)

```

```

await websocket.accept()

Subscribe to project events
event_subscription = await event_bus.subscribe(
 f"project.{project_id}.*", websocket
)

try:
 while True:
 # Keep connection alive and handle incoming messages
 data = await websocket.receive_text()
 message = json.loads(data)

 if message["type"] == "ping":
 await websocket.send_json({"type": "pong"})
 elif message["type"] == "subscribe":
 await handle_subscription_request(message, websocket)

except WebSocketDisconnect:
 await event_subscription.unsubscribe()

```

## 2. GraphQL API for Complex Queries



python

```

import strawberry
from strawberry.fastapi import GraphQLRouter

@strawberry.type
class Project:
 id: str
 name: str
 status: str
 agents: List['Agent']
 tasks: List['Task']
 metrics: 'ProjectMetrics'

@strawberry.type
class Agent:
 id: str
 type: str
 specialization: str
 status: str
 current_tasks: List['Task']
 performance_metrics: 'AgentMetrics'

@strawberry.type
class Task:
 id: str
 title: str
 description: str
 status: str
 assigned_agent: Agent
 estimated_effort: float
 actual_effort: Optional[float]
 quality_score: Optional[float]

@strawberry.type
class Query:
 @strawberry.field
 async def project(self, id: str, info: Info) -> Optional[Project]:
 """Get detailed project information with nested data"""

 user = await get_current_user_from_context(info.context)
 await check_project_access(id, user)

 # Use DataLoader pattern for efficient data fetching
 project_loader = info.context["project_loader"]

```

```
project = await project_loader.load(id)
```

```
return project
```

```
@strawberry.field
```

```
async def projects(
```

```
 self,
```

```
 filter: Optional[ProjectFilter] = None,
```

```
 pagination: Optional[PaginationInput] = None,
```

```
 info: Info = None
```

```
) -> ProjectConnection:
```

```
 """Get filtered and paginated list of projects"""
```

```
 user = await get_current_user_from_context(info.context)
```

```
 # Build query with filters
```

```
 query_builder = ProjectQueryBuilder()
```

```
 query = query_builder.build_query(filter, user.organization_id)
```

```
 # Execute with pagination
```

```
 results = await project_service.get_projects(
```

```
 query, pagination or PaginationInput()
```

```
)
```

```
 return ProjectConnection(
```

```
 edges=[ProjectEdge(node=project) for project in results.projects],
```

```
 page_info=results.page_info,
```

```
 total_count=results.total_count
```

```
)
```

```
@strawberry.type
```

```
class Mutation:
```

```
 @strawberry.mutation
```

```
 async def create_project(
```

```
 self,
```

```
 input: CreateProjectInput,
```

```
 info: Info
```

```
) -> CreateProjectPayload:
```

```
 """Create a new project with intelligent agent assignment"""
```

```
 user = await get_current_user_from_context(info.context)
```

```
 try:
```

```
 # Create project with auto-agent assignment
```

```

 project = await project_service.create_project_with_agents(
 input, user.organization_id
)

 return CreateProjectPayload(
 project=project,
 success=True,
 errors=[]
)
 except ValidationError as e:
 return CreateProjectPayload(
 project=None,
 success=False,
 errors=e.errors
)

Create GraphQL schema
schema = strawberry.Schema(query=Query, mutation=Mutation)
graphql_app = GraphQLRouter(schema)

Add to FastAPI app
app.include_router(graphql_app, prefix="/graphql")

```

---

## DISASTER RECOVERY & BACKUP

### 1. Multi-Region Backup Strategy

python

```

class DisasterRecoveryManager:
 """Comprehensive disaster recovery and backup management"""

 def __init__(self):
 self.backup_scheduler = BackupScheduler()
 self.replication_manager = ReplicationManager()
 self.recovery_orchestrator = RecoveryOrchestrator()

 async def setup_multi_region_backup(self, customer_id: str):
 """Set up multi-region backup for customer data"""

 backup_config = BackupConfiguration(
 customer_id=customer_id,
 backup_frequency="hourly",
 retention_policy={
 "hourly": "24h",
 "daily": "30d",
 "weekly": "12w",
 "monthly": "12m"
 },
 regions=["us-east-1", "us-west-2", "eu-west-1"],
 encryption_enabled=True,
 compression_enabled=True
)

 # Set up automated backups
 await self.backup_scheduler.schedule_backups(backup_config)

 # Configure cross-region replication
 await self.replication_manager.setup_replication(backup_config)

 # Test backup integrity
 await self.test_backup_integrity(customer_id)

 async def execute_disaster_recovery(self, incident: DisasterIncident):
 """Execute disaster recovery procedures"""

 recovery_plan = await self.generate_recovery_plan(incident)

 # Notify stakeholders
 await self.notify_incident_response_team(incident, recovery_plan)

 # Execute recovery steps

```

```

for step in recovery_plan.steps:
 step_result = await self.execute_recovery_step(step)

 if not step_result.success:
 await self.escalate_recovery_failure(step, step_result)
 break

 await self.update_recovery_status(incident.id, step.id, step_result)

Validate recovery
validation_result = await self.validate_recovery(incident)

return RecoveryResult(
 incident=incident,
 recovery_plan=recovery_plan,
 validation_result=validation_result,
 recovery_time=datetime.utcnow() - incident.started_at
)

class BackupManager:
 """Manages all backup operations across the platform"""

 async def backup_agent_state(self, agent_id: str):
 """Create comprehensive backup of agent state and knowledge"""

 # Backup agent configuration
 agent_config = await self.agent_registry.get_agent_config(agent_id)

 # Backup agent memory and context
 agent_memory = await self.memory_manager.export_agent_memory(agent_id)

 # Backup learned patterns and optimizations
 agent_knowledge = await self.knowledge_manager.export_agent_knowledge(agent_id)

 # Create backup artifact
 backup_artifact = AgentBackup(
 agent_id=agent_id,
 timestamp=datetime.utcnow(),
 config=agent_config,
 memory=agent_memory,
 knowledge=agent_knowledge,
 version=await self.get_agent_version(agent_id)
)

```

```

 # Store in multiple locations
 backup_locations = await self.store_backup_artifact(backup_artifact)

 return backup_locations

 async def restore_agent_from_backup(
 self,
 agent_id: str,
 backup_timestamp: datetime
):
 """Restore agent from backup to specific point in time"""

 # Find appropriate backup
 backup_artifact = await self.find_backup_artifact(
 agent_id, backup_timestamp
)

 if not backup_artifact:
 raise BackupNotFoundError(
 f"No backup found for agent {agent_id} at {backup_timestamp}"
)

 # Stop current agent
 await self.agent_manager.graceful_shutdown_agent(agent_id)

 # Restore from backup
 restored_agent = await self.restore_agent_from_artifact(backup_artifact)

 # Validate restoration
 validation_result = await self.validate_agent_restoration(restored_agent)

 if validation_result.is_valid:
 # Restart agent
 await self.agent_manager.start_agent(restored_agent)
 return restored_agent
 else:
 raise RestorationError(
 f"Agent restoration validation failed: {validation_result.errors}"
)

```

## 2. High Availability Architecture



yaml

# *# Kubernetes High Availability Configuration*

apiVersion: v1

kind: ConfigMap

metadata:

name: aamas-ha-config

data:

ha-config.yaml: |

high\_availability:

enabled: true

replication\_factor: 3

auto\_failover: true

health\_check\_interval: 30s

load\_balancing:

strategy: "round\_robin"

health\_checks:

- endpoint: "/health"

interval: 10s

timeout: 5s

failure\_threshold: 3

data\_replication:

mode: "synchronous"

regions: ["us-east-1", "us-west-2", "eu-west-1"]

consistency\_level: "strong"

backup\_schedule:

frequency: "0 \*/4 \* \* \*" # Every 4 hours

retention: "30d"

verification: true

---

apiVersion: apps/v1

kind: StatefulSet

metadata:

name: aamas-database-cluster

spec:

serviceName: aamas-db-service

replicas: 3

selector:

matchLabels:

app: aamas-database

template:

```
metadata:
 labels:
 app: aamas-database
spec:
 containers:
 - name: postgresql
 image: postgres:14-alpine
 env:
 - name: POSTGRES_DB
 value: "aamas"
 - name: POSTGRES_USER
 valueFrom:
 secretKeyRef:
 name: db-credentials
 key: username
 - name: POSTGRES_PASSWORD
 valueFrom:
 secretKeyRef:
 name: db-credentials
 key: password
 - name: POSTGRES_REPLICATION_MODE
 value: "master"
 volumeMounts:
 - name: postgres-storage
 mountPath: /var/lib/postgresql/data
 ports:
 - containerPort: 5432
 livenessProbe:
 exec:
 command:
 - pg_isready
 - -U
 - $(POSTGRES_USER)
 initialDelaySeconds: 30
 periodSeconds: 10
 readinessProbe:
 exec:
 command:
 - pg_isready
 - -U
 - $(POSTGRES_USER)
 initialDelaySeconds: 5
 periodSeconds: 5
 volumeClaimTemplates:
```

```
- metadata:
 name: postgres-storage
spec:
 accessModes: ["ReadWriteOnce"]
 resources:
 requests:
 storage: 100Gi
 storageClassName: "fast-ssd"
```

---

## COST OPTIMIZATION

### 1. Intelligent Resource Management

python

```

class CostOptimizationEngine:
 """AI-powered cost optimization for the platform"""

 def __init__(self):
 self.usage_analyzer = UsageAnalyzer()
 self.cost_predictor = CostPredictor()
 self.resource_optimizer = ResourceOptimizer()

 async def optimize_customer_costs(self, customer_id: str):
 """Analyze and optimize costs for a specific customer"""

 # Analyze current usage patterns
 usage_analysis = await self.usage_analyzer.analyze_customer_usage(
 customer_id, timeframe="30d"
)

 # Predict future costs
 cost_prediction = await self.cost_predictor.predict_costs(
 customer_id, usage_analysis
)

 # Generate optimization recommendations
 optimizations = await self.generate_cost_optimizations(
 usage_analysis, cost_prediction
)

 # Auto-apply safe optimizations
 applied_optimizations = []
 for optimization in optimizations:
 if optimization.risk_level == "low" and optimization.auto_apply:
 result = await self.apply_optimization(customer_id, optimization)
 applied_optimizations.append(result)

 return CostOptimizationResult(
 current_costs=usage_analysis.total_cost,
 predicted_costs=cost_prediction.monthly_cost,
 potential_savings=sum(opt.estimated_savings for opt in optimizations),
 applied_optimizations=applied_optimizations,
 recommendations=optimizations
)

 async def dynamic_resource_scaling(self):
 """Continuously optimize resource allocation across the platform"""

```

```

while True:
 # Get current resource utilization
 utilization = await self.get_platform_utilization()

 # Identify optimization opportunities
 opportunities = await self.identify_scaling_opportunities(utilization)

 for opportunity in opportunities:
 if opportunity.action == "scale_down":
 await self.scale_down_resources(opportunity)
 elif opportunity.action == "consolidate":
 await self consolidate_workloads(opportunity)
 elif opportunity.action == "spot_instance":
 await self.migrate_to_spot_instances(opportunity)

 # Sleep for 5 minutes before next check
 await asyncio.sleep(300)

```

```

class ResourceOptimizer:

```

```

 """Optimizes compute resources for AI agents"""

```

```

 async def optimize_agent_resources(self, agent_id: str):

```

```

 """Optimize compute resources for a specific agent"""

```

```

 # Analyze agent performance history

```

```

 performance_history = await self.get_agent_performance_history(
 agent_id, timeframe="7d"
)

```

```

 # Determine optimal resource allocation

```

```

 optimal_config = await self.calculate_optimal_resources(
 agent_id, performance_history
)

```

```

 # Check if changes are beneficial

```

```

 cost_benefit = await self.analyze_cost_benefit(
 agent_id, optimal_config
)

```

```

 if cost_benefit.net_benefit > 0:

```

```

 # Apply resource changes

```

```

 await self.apply_resource_changes(agent_id, optimal_config)

```

```
 return ResourceOptimizationResult(
 agent_id=agent_id,
 old_config=performance_history.current_config,
 new_config=optimal_config,
 estimated_savings=cost_benefit.monthly_savings,
 performance_impact=cost_benefit.performance_impact
)

 return None # No beneficial changes found
```

## 2. Cost Monitoring & Alerting



python

```

class CostMonitoringSystem:
 """Real-time cost monitoring and alerting"""

 def __init__(self):
 self.cost_tracker = CostTracker()
 self.budget_manager = BudgetManager()
 self.alert_manager = AlertManager()

 async def setup_customer_cost_monitoring(self, customer_id: str):
 """Set up comprehensive cost monitoring for a customer"""

 # Set up cost tracking
 await self.cost_tracker.setup_customer_tracking(customer_id)

 # Configure budget alerts
 budget_config = await self.budget_manager.get_customer_budget(customer_id)

 alerts = [
 CostAlert(
 name=f"budget_50_percent_{customer_id}",
 condition=f"monthly_spend > {budget_config.monthly_limit * 0.5}",
 severity="info",
 message="You've reached 50% of your monthly budget"
),
 CostAlert(
 name=f"budget_80_percent_{customer_id}",
 condition=f"monthly_spend > {budget_config.monthly_limit * 0.8}",
 severity="warning",
 message="You've reached 80% of your monthly budget"
),
 CostAlert(
 name=f"budget_exceeded_{customer_id}",
 condition=f"monthly_spend > {budget_config.monthly_limit}",
 severity="critical",
 message="Monthly budget exceeded!"
)
]

 for alert in alerts:
 await self.alert_manager.register_alert(alert)

 async def generate_cost_report(self, customer_id: str, period: str):
 """Generate detailed cost breakdown report"""

```

```
cost_data = await self.cost_tracker.get_cost_breakdown(
 customer_id, period
)

Analyze cost trends
trends = await self.analyze_cost_trends(cost_data)

Generate optimization suggestions
suggestions = await self.generate_cost_suggestions(cost_data, trends)

return CostReport(
 customer_id=customer_id,
 period=period,
 total_cost=cost_data.total,
 breakdown=cost_data.breakdown,
 trends=trends,
 optimization_suggestions=suggestions,
 generated_at=datetime.utcnow()
)
```

---

## COMPLIANCE & GOVERNANCE

### 1. Data Governance Framework

python

```

class DataGovernanceManager:
 """Manages data governance, privacy, and compliance"""

 def __init__(self):
 self.privacy_manager = PrivacyManager()
 self.compliance_checker = ComplianceChecker()
 self.audit_logger = AuditLogger()
 self.data_classifier = DataClassifier()

 async def setup_customer_governance(self, customer_id: str, requirements: ComplianceRequirements):
 """Set up data governance for a customer based on their compliance needs"""

 # Classify customer data sensitivity
 data_classification = await self.data_classifier.classify_customer_data(
 customer_id
)

 # Set up appropriate data handling policies
 policies = await self.generate_data_policies(
 data_classification, requirements
)

 # Configure privacy controls
 await self.privacy_manager.setup_privacy_controls(
 customer_id, policies.privacy_policy
)

 # Set up compliance monitoring
 await self.compliance_checker.setup_compliance_monitoring(
 customer_id, requirements, policies
)

 # Configure audit logging
 await self.audit_logger.setup_customer_auditing(
 customer_id, policies.audit_policy
)

 return GovernanceSetup(
 customer_id=customer_id,
 data_classification=data_classification,
 policies=policies,
 compliance_requirements=requirements
)

```

```

async def handle_data_deletion_request(self, customer_id: str, deletion_request:
 """Handle GDPR/CCPA data deletion requests"""

 # Validate deletion request
 validation = await self.validate_deletion_request(deletion_request)
 if not validation.is_valid:
 raise InvalidDeletionRequest(validation.errors)

 # Find all customer data
 data_inventory = await self.create_data_inventory(
 customer_id, deletion_request.scope
)

 # Execute deletion across all systems
 deletion_results = []
 for data_location in data_inventory.locations:
 result = await self.delete_data_from_location(
 data_location, deletion_request
)
 deletion_results.append(result)

 # Verify deletion completeness
 verification = await self.verify_data_deletion(
 customer_id, deletion_request, deletion_results
)

 # Log deletion for compliance
 await self.audit_logger.log_data_deletion(
 customer_id=customer_id,
 deletion_request=deletion_request,
 results=deletion_results,
 verification=verification
)

 return DataDeletionResult(
 request=deletion_request,
 results=deletion_results,
 verification=verification,
 completed_at=datetime.utcnow()
)

```

```

class ComplianceFramework:

```

```

 """Comprehensive compliance framework supporting multiple standards"""

```

```

supported_standards = [
 "SOC2_TYPE2",
 "ISO27001",
 "GDPR",
 "CCPA",
 "HIPAA",
 "PCI_DSS"
]

async def assess_compliance_posture(self, customer_id: str):
 """Assess current compliance posture across all applicable standards"""

 customer_profile = await self.get_customer_compliance_profile(customer_id)

 assessments = {}
 for standard in customer_profile.applicable_standards:
 assessment = await self.assess_standard_compliance(
 customer_id, standard
)
 assessments[standard] = assessment

 # Generate overall compliance score
 overall_score = self.calculate_overall_compliance_score(assessments)

 # Identify critical gaps
 critical_gaps = self.identify_critical_gaps(assessments)

 # Generate remediation plan
 remediation_plan = await self.generate_remediation_plan(critical_gaps)

 return ComplianceAssessment(
 customer_id=customer_id,
 overall_score=overall_score,
 standard_assessments=assessments,
 critical_gaps=critical_gaps,
 remediation_plan=remediation_plan,
 next_assessment_date=datetime.utcnow() + timedelta(days=90)
)

```

---

## CONCLUSION & NEXT STEPS

This comprehensive system architecture provides AAMAS with:

### ✓ **Technical Foundation**

- **Scalable microservices architecture** supporting multi-tenant SaaS operations
- **Intelligent agent orchestration** with sophisticated communication protocols
- **Enterprise-grade security** with zero-trust principles and comprehensive auditing
- **High availability & disaster recovery** with multi-region redundancy

### ✓ **Operational Excellence**

- **Real-time monitoring & observability** with predictive alerting
- **Cost optimization** through AI-powered resource management
- **Compliance framework** supporting major industry standards
- **Developer-friendly APIs** with both REST and GraphQL interfaces

### ✓ **Business Enablement**

- **Multi-tier pricing** support with granular resource allocation
- **Customer analytics** with actionable insights and ROI tracking
- **Integration ecosystem** supporting 50+ development tools
- **White-label capabilities** for enterprise customization

### **Implementation Roadmap**

#### **Phase 1 (Months 1-3): Core Infrastructure**

- Deploy basic microservices architecture
- Implement agent communication framework
- Set up development/staging environments
- Basic security and monitoring

#### **Phase 2 (Months 4-6): Agent Intelligence**

- Deploy specialized agent implementations
- Integrate AI models and knowledge systems
- Implement predictive analytics
- Customer onboarding system

#### **Phase 3 (Months 7-9): Enterprise Features**



- Advanced security and compliance
- Multi-region deployment
- Cost optimization engine
- Advanced analytics and reporting

**Phase 4 (Months 10-12): Scale & Optimize**

- Performance optimization
- Advanced integrations
- Global expansion features
- AI model fine-tuning

This architecture positions AAMAS as a truly enterprise-ready platform capable of transforming how software development teams operate, with the intelligence and automation needed to achieve the ambitious business goals outlined in your proposal.